

Field2VC: Generating Verification Criteria for Automotive SPICE SYS.2 from Field-Failure Evidence

Anonymous Author(s)

Submission under double-blind review

Affiliation withheld for anonymity

Abstract—In Automotive SPICE, SYS.2 system requirements must satisfy verifiability, with verification criteria (VCs) inherent in the requirements text; for safety-relevant ECUs, ISO 26262 [1] reinforces this obligation. In practice, VCs are often incomplete in a recurring way: engineers rarely foresee failure modes they have not previously encountered. We present Field2VC, which grounds LLM-based VC generation in two complementary failure sources—developer-reported defects on GitHub, which expose code-level boundary conditions absent from specifications, and consumer complaints in the NHTSA ODI database, which expose vehicle-level failures engineers do not anticipate. Field2VC filters each report into a structured failure representation, retrieves the failures most relevant to a requirement, and generates specific, testable VCs. In an evaluation of 489 generated VCs over a 350-requirement industrial benchmark—each VC independently scored by two of three external raters, with an internal adjudicator resolving splits—failure grounding raises the conditional acceptance rate among evaluated VCs from a 45.3 % no-failure-evidence baseline to 65.5 %–67.2 % (odds ratio 2.3–2.5 $\times$; Holm-corrected $p < 0.001$). An adjudication-sensitivity analysis shows this gain is uneven: on the agreed subset, all Channel A items were accepted without E1 adjudication (37/37 versus 119/128 for the baseline), whereas the NHTSA channel fell to 20/39 (51.3 %). The results suggest that a failure source’s usefulness depends on vocabulary alignment: each source is effective on the requirement subdomain whose language its filtered failures match—developer issues for the communication and diagnostics stack, regulator complaints for body-control. We provide the benchmark, all generated VCs, and per-rater labels.

Index Terms—Automotive software engineering, Automotive SPICE, failure report mining, field-failure evidence, large language models (LLMs), requirements engineering, system requirements analysis, verification criteria

I. INTRODUCTION

In automotive software engineering, the cost of repairing a defect rises sharply the later it is found; classic, if debated, estimates put a post-delivery fix at up to 100 times the cost of a requirements-phase fix [2] (Fig. 1). The obligation, however, does not rest on any particular cost multiplier. ASPICE SYS.2 requires system requirements to satisfy *verifiability*: verification criteria should be inherent in the requirements text [3]. For safety-relevant ECUs this obligation is reinforced by ISO 26262, whose functional-safety lifecycle requires each requirement to be verified before integration [1].

Writing thorough VCs, however, often fails in a specific and repeatable way. A requirement specifies intended behavior

under nominal conditions; a complete VC must additionally probe boundary conditions and fault modes. Reviewers reason forward from the requirement text, so the criteria they write can only exercise behaviors the specification already anticipates, while field failures, almost by definition, occur where it did not. The blind spot is structural rather than a matter of effort: a VC that omits an unforeseen edge case passes internal review precisely because no reviewer holds the evidence that would expose it (Fig. 1).

Closing this gap requires failure evidence from outside the project, matched to the abstraction level at which a SYS.2 requirement is verified—the *system* level, not the software unit. Defect trackers are a natural candidate, mined by recent work such as CrUISE-AC [4] to enrich agile acceptance criteria: developer-reported defects describe faults as engineers observe them—isolated to a module, reproducible, often with a root cause. But failures emerging from cross-subsystem interactions and months of customer usage rarely surface in any repository, because no developer observes the deployed fleet. Regulator-collected complaints capture exactly this stratum: NHTSA’s ODI database records faults as the vehicle exhibits them, at fleet scale, under conditions no test plan enumerates [5], [6]. The two corpora therefore appear complementary under the evaluated mapping; our evaluation supports this split empirically—complaint-grounded VCs are accepted more often for body and chassis requirements, defect-grounded VCs for communication and diagnostics (§V-B). We use NHTSA as a public, reproducible proxy for the warranty and field-return data an OEM holds internally but cannot release; the pipeline is agnostic to the vehicle-level channel’s source. Manual inspection of 200 sampled requirement–evidence pairs indicates this mapping stage reaches **73.5 %** precision (§IV).

This motivates our central question: given a system requirement, can generation grounded in cross-level field-failure evidence produce verification criteria that practicing engineers would accept into a SYS.2 baseline? We treat expert acceptance as the measurable proxy for usefulness, since no oracle for “complete” VC coverage exists.

This paper presents **Field2VC**, an automated VC generation mechanism grounded in real failure evidence. Field2VC extracts defect context from two channels—GitHub (code-level failures) and NHTSA (vehicle-level failures)—filters domain-specific noise from the raw records, and uses the retained




Fig. 1. ASPICE V-model cost escalation. A missing VC that passes SYS.2 review surfaces only at system-level testing or as a field complaint, where re-verification is far costlier. Field2VC shifts detection left by grounding VC generation in field-failure evidence at SYS.2 (the *Field2VC scope* annotation); the green arc is the evidence feedback loop from field operation back to SYS.2. SWE-layer processes (SWE.1–SWE.6) are out of scope.

context as a semantic constraint so a large language model (LLM) generates verifiable VCs covering historical failure scenarios.

a) Contributions:

- **An empirical finding:** a four-expert evaluation of 489 generated VCs (307 unique industrial requirements) shows that field-failure grounding raises expert acceptance over a no-failure-evidence baseline, and suggests a design heuristic—a failure source appears useful where its filtered vocabulary aligns with the target requirement subdomain (§V).
- **An automated, evidence-grounded pipeline** combining developer-reported defects (GitHub) and regulator-collected complaints (NHTSA) at distinct abstraction levels, with an abstention mechanism that declines when retrieved evidence is insufficient (§III).
- **An industrial benchmark:** 350 de-identified SYS.2 requirements (307 evaluated) with expert-annotated golden VCs, all generated VCs, per-rater labels, and analysis scripts (Data Availability).

Our results extend CrUISE-AC’s premise that external issue knowledge can enrich requirements artifacts [4], but move it to a safety-critical setting: the evidence is drawn from two channels at *distinct abstraction levels*, and the target artifact is an ASPICE-mandated verification criterion for a safety-relevant ECU rather than an agile acceptance criterion.

b) *Data Availability:* The de-identified anonymous replication package is available for review at <https://anonymous.4open.science/r/apsec2026-artifact-C562/>. It contains the 350-item SYRS corpus with golden VCs, all 489 generated VCs

with per-rater and final labels, the adjudication record, the evaluation protocol, prompts, model identifiers, raw model outputs, pipeline scripts, and statistical reproducers. Proprietary signal names are masked by stable placeholders; the redaction script is included.

II. RELATED WORK

A. Automated Requirements Completion within Specifications

Requirements completeness research falls into two streams. *Refinement* approaches improve individual requirements through semantic enrichment without adding new external data, using techniques such as NL-level templating and domain-knowledge-driven pattern systems, e.g., the EARS controlled requirements syntax [7]. *Derivation* approaches generate new requirements by exploiting relationships within or between existing requirement items.

LLM-based derivation has advanced rapidly [8]. F2SRD [9] retrieves the applicable OWASP ASVS verification requirements for each functional specification and guides GPT-4 to derive security requirements grounded in them. Such approaches draw their completion signal from within-specification logic or applicable standards, not from empirical failure evidence; our work is complementary: we exploit *external* failure repositories to surface conditions that no single specification can reveal.

Closest in spirit is TVR [10], which applies retrieval-augmented generation (RAG) to validate and recover missing traceability links between stakeholder and system requirements in industrial automotive projects. TVR targets *traceability* (linking existing requirements) rather than *VC completion*

(generating new complementary criteria); its retrieval corpus is the specification itself, whereas ours is external failure evidence. These two directions show that RAG-based requirements engineering is being applied across multiple automotive sub-tasks [11], [12].

B. Enriching Requirements from External Repositories

Mining bug reports and issue trackers for RE has a long history [13], [14]. Do et al. [15] extract requirements from similar products via doc2vec and BIRCH clustering; Almonte et al. [16] prompt LLMs to derive quality-attribute NFRs from functional specifications; Koscinski et al. [17] use RelGAN to transform functional information into security requirements. These mine requirements-like data from whatever corpora are available; we instead mine two repositories from *different domains*—a defect tracker and a regulator-operated complaint database—and show that the knowledge source is a central design variable for VC quality.

CrUISE-AC [4] is the leading crowd-data-driven system over external repositories, mining public issue trackers (seven e-commerce systems, >165 k raw issue records, plus two content-management systems) for issues matching Gherkin acceptance criteria, achieving 82% approval in e-commerce under a four-expert (≥ 3 -of-4) majority vote with Gwet’s AC1=0.74. Wang et al. [18] likewise generate acceptance criteria from multi-modal requirements with LLMs, but target user-facing rather than safety-critical requirements.

Field2VC differs from CrUISE-AC in three respects. First, our domain is the automotive system-requirements level (AS-PICE SYS.2) for safety-relevant ECUs. Second, we evaluate against a three-criterion acceptance protocol (Completeness ≥ 3 , Novelty=Y, Usefulness=Y) rather than a binary judgment, making direct percentage comparisons misleading. Third, our ablation contrasts LLM classification and embedding matching as grounding mechanisms, primarily exposing the effect of the filtering-enabled abstention decision rather than the matching method alone (§V-C).

C. LLM-based Requirements Engineering

Recent work applies LLMs to diverse software engineering tasks [19], [20]: MARE [21] proposes a multi-agent collaboration framework for requirements elicitation. Like the others, it targets neither VC completion nor the effect of grounding in external failure evidence—the key contribution of Field2VC. In the automotive-safety domain specifically, LLMs have been applied to code generation and verification [22], [23], functional-safety and security design [24], and RAG-based derivation of safety requirements from standards [25]; these operate on code or upstream safety requirements rather than completing verification criteria for existing SYS.2 specifications.

III. FIELD2VC: TWO-CHANNEL VC COMPLETION PIPELINE

Field2VC is a five-stage pipeline that transforms raw failure records from two complementary external channels into

candidate missing Verification Criteria (VCs) for automotive ECU SYRS items, each traceable to one or more source records (Fig. 2). We illustrate it throughout with a real Unified Diagnostic Services (UDS, ISO 14229) Security Access requirement (SYRS.1018_A) and its GitHub field failure (python-udsoncan#51)—a Channel A trace that surfaces protocol-transport boundary conditions absent from specification. Table I complements it with a Channel B (NHTSA) trace, where complaint evidence surfaces a real-world electromechanical scenario the golden VC did not anticipate.

TABLE I
WORKED TRACE: AN ACCEPTED CHANNEL B (NHTSA) VC TRACED THROUGH ALL FIVE STAGES. FAILURE TEXT IS VERBATIM FROM THE PUBLIC SOURCE; SYRS AND GOLDEN VC ARE PARAPHRASED FROM THE DE-IDENTIFIED CORPUS.

Channel B: NHTSA, SYRS.0177 (low-beam robustness under AC flicker). SYRS: do <i>not</i> switch off low-beam headlights when exposed to synthetic AC-flicker light sources if ambient brightness remains <1000 lux for more than 4 s. The golden VC tests the <i>nominal AC-flicker path only</i>. <i>Matched failure:</i> ODI#11557628 (sim. 0.60; 2022 Toyota Camry): “low beam headlights flashing when high beams are on... also flash on ignition.” <i>Triple:</i> (high-beam activation or ignition key-on low-beam LEDs remain steadily illuminated bilateral intermittent flashing (right front more severe); state-stuck fault). <i>Generated VC:</i> fault-injection in an AC-flicker environment (<1000 lux sustained >4 s): activate high beams and separately cycle the ignition, verifying that low beams neither flash nor extinguish in either case—electromechanical interaction scenarios absent from the golden VC. Accepted (E3: C=3; E4: C=3; N=Y, U=Y; unanimous).
--

A. Stage 1: Issue Collection

We define two channels providing complementary failure evidence:

Channel A: GitHub Issues. We index issues from four open-source repositories covering ECU-adjacent stacks: pylessard/python-udsoncan (UDS diagnostics), python-can-isotp (ISO-TP transport), hardbyte/python-can (CAN bus), and ecubus/EcuBus-Pro (ECU test toolchain). We retrieved **696 issues** via the GitHub REST API,¹ prioritizing closed issues labeled {bug, defect, regression, safety} and supplementing with unlabeled closed issues for coverage. The running example fetches python-udsoncan#51: “SecurityAccess does not handle already-unlocked server.”

Channel B: NHTSA Complaints. We query the NHTSA ODI consumer-complaint database [26] using component keywords matching automotive ECU subsystems. We retrieved **364 complaint records**, each containing a free-text consumer narrative plus structured vehicle and component metadata. Unlike curated recall notices, complaint narratives are raw, colloquial, and noisy field reports, which is why the Stage-2 filter matters most for this channel.

¹<https://docs.github.com/en/rest/issues>




Fig. 2. Field2VC five-stage pipeline. Failure records from GitHub (Ch-A) and NHTSA (Ch-B) are filtered into $\langle \text{Trigger} | \text{Protocol Behavior} | \text{Fault Mode} \rangle$ triples (Stage 2, Qwen-Plus), embedding-matched to each SYRS item via cosine similarity (Stage 3, `text-embedding-v3`), and used to prompt VC generation (Stage 4, Qwen-Max). The generator is never forced to produce a VC: it may decline by emitting the fixed marker `NO_NOVEL_VC_FOUND` when the evidence reveals no genuinely missing condition, the pipeline’s post-retrieval relevance gate. Four domain experts review the surviving VCs under the three-criterion acceptance protocol (Stage 5). The diagram traces the 215 channel VCs (Ch-A + Ch-B); the full study adds 274 no-failure-evidence baseline VCs for 489 total (Table II).

B. Stage 2: Issue Decomposition

We use an LLM (Qwen-Plus) to filter each raw record into a structured triple:

$\langle \text{Trigger} | \text{Protocol Behavior} | \text{Fault Mode} \rangle$

Trigger is the activating condition, *Protocol Behavior* the expected ECU action per the relevant standard, and *Fault Mode* the observed deviation; filtering reduces vocabulary mismatch and suppresses implementation-layer detail irrelevant to SYS.2. The core instruction is “*Extract the technical essence from this report*”, with a mandatory escape route: records not about automotive system behavior (e.g., Python packaging, documentation, consumer venting) return all three fields as `null`² and are dropped—the point where off-topic noise leaves the pipeline. The running example, issue #51, reduces to $\langle 0x27 \text{ requestSeed while already unlocked} | \text{client detects all-zero seed and skips sendKey} | \text{key derived from zero seed} \rightarrow \text{Negative Response Code (NRC) } 0x24 \text{ rejection} \rangle$.

C. Stage 3: Context Retrieval

For each SYRS item s and filtered triple t_i we compute cosine similarity [27], [28] using `text-embedding-v3` (Alibaba DashScope; snapshot April–May 2026), retrieving the top- $K = 5$ triples as the *failure context* $\mathcal{F}(s) = \{t_1, \dots, t_5\}$; $K=5$ is a pragmatic trade-off between computational cost and retrieval diversity. We impose *no* hard similarity cut-off: every item proceeds with its top-5 context and the relevance decision is delegated to generation [29], where the model abstains (`NO_NOVEL_VC_FOUND`) if the evidence reveals

²All Stage-2 and Stage-4 prompts are released verbatim in the replication package (prompts/).

no genuinely missing condition (§III-D). This mirrors, post-retrieval, the pre-retrieval LLM relevance classifier of CrUISE-AC [4]—a choice our ablation examines directly (RQ3, §V-C). For SYRS.1018_A, issue #51 ranks second at cosine 0.69 and enters the top-5 context. Embedding matching has two advantages over LLM binary classification: it scales to thousands of pairs without per-pair LLM calls, and it yields a continuous, auditable relevance score released with every match.

D. Stage 4: VC Generation

We prompt Qwen-Max with three inputs: the SYRS item, its existing golden VCs (as *do-not-repeat* context operationalising “missing”), and the failure context $\mathcal{F}(s)$. The prompt requires VCs that (a) genuinely test the requirement (not the protocol library the failure was reported against), (b) are not redundant with the golden VCs, and (c) are inspired by a real failure pattern in $\mathcal{F}(s)$ and testable at ECU system-integration level. It closes with explicit permission to decline—“*if no genuinely useful missing VC can be identified, output NO_NOVEL_VC_FOUND*” (the *sentinel*). Declining is expected, not an error: rather than hallucinate a plausible criterion [30], the pipeline abstains when the evidence is irrelevant. For the running example, the model turns issue #51’s triple into a negative test of the already-unlocked path ($0x27$ at L1 unlocked, send $0x31 \ 01 \ 10 \ 08$, expect no NRC $0x33$ securityAccessDenied)—a scenario neither the SYRS nor its golden VC exercises. The prompt uses a SYS.5 test-design persona, but the output is evaluated as a SYS.2 verification criterion coexisting with the requirement.

E. Stage 5: Expert Evaluation

Each generated VC is independently scored by at least two of the three external raters (E2, E3, E4) on three dimensions under a **three-criterion acceptance protocol**:

- **C (Completeness)**: Score 1–5. Does the SYRS item omit this VC? ($C \geq 3$ is the acceptance threshold.)
- **N (Novelty)**: Y/N. Is this VC absent from the existing SYRS item and its associated VCs?
- **U (Usefulness)**: Y/N. Is this VC actionable for downstream test design and software qualification?

A VC is accepted if and only if $C \geq 3$ AND $N = Y$ AND $U = Y$. For the main evaluation, each VC receives two external ratings; when those ratings split, E1 resolves the final label under the same protocol. The running-example VC was accepted ($C=4$, $N=Y$, $U=Y$) after E1 adjudicated a novelty split between E2 and E3.

F. Ablation Variants

Two ablation conditions probe filtering and matching on the same 50-item subset. **Abl-NF (No Filter)** skips Stage 2, embedding the raw issue title + body (first 500 characters) and matching to SYRS by cosine similarity. **Abl-LM (LLM-Match)** also skips Stage 2, but replaces embedding similarity with an LLM binary relevance classifier (Qwen-Plus), approximating CrUISE-AC’s LLM matching family without its multi-LLM ensemble [4]. Since the filter + LLM-match cell is absent, RQ3 is an ablation study rather than a full factorial decomposition.

IV. EXPERIMENT DESIGN

A. Research Questions

- **RQ1 (Effectiveness)**: Does failure-grounded prompting generate VCs with significantly higher expert acceptance than a no-failure-evidence baseline?
- **RQ2 (Channel-Domain Fit)**: The two channels are mapped to disjoint requirement subdomains by construction. Within each subdomain (body-control/chassis, communication stack/state management, diagnostics), how effective is the assigned channel relative to the no-failure-evidence baseline, and how do the two channels divide labour across subdomains?
- **RQ3 (Ablation)**: What is the contribution of (a) failure-triple filtering and (b) embedding-based matching to VC quality?
- **RQ4 (Sentinel-Adherence Sensitivity)**: How consistently do different LLMs honor the generation-refusal sentinel used by the pipeline?

B. Dataset

a) *SYRS corpus*: We apply the pipeline to 350 unique automotive ECU SYRS items drawn from a proprietary corpus of 675: the 350 are every item falling in the two channels’ target subdomains, with items in other domains (e.g., powertrain, ADAS, gateway) out of scope (§VI). The channel mapping is disjoint by construction: Ch-A = {communication stack (UDS,

CAN), diagnostics} (208 items) and Ch-B = {body control (steering-column lock, parking brake, steering-column control module), IO/HMI, functional safety, system-state management} (142 items). This canonical mapping governs every per-category breakdown (Table IV). All items follow the ASPICE SYS.2 format (explicit actor, trigger, quantitative boundary, expected behavior) and carry human-annotated golden VCs.

b) *Generated VCs*: Across both channels and the no-failure-evidence baseline, 489 VCs entered expert evaluation (Table II):

TABLE II
DATASET AND PER-CONDITION VC FLOW. CH-A AND CH-B COVER DISJOINT SYRS SUBSETS (208 + 142 = 350) AND ABSTAIN VIA THE NO_NOVEL_VC_FOUND SENTINEL. “ITEMS”/“ABST.” ARE COUNTED PER SYRS ITEM WHILE “CAND.”/“EVAL.” ARE PER VC, AS A NON-ABSTAINING ITEM MAY EMIT MORE THAN ONE CANDIDATE; THE CAND. → EVAL. GAP THEN REMOVES DUPLICATE/MALFORMED CANDIDATES. SOURCES: 696 GITHUB ISSUES, 364 NHTSA COMPLAINTS.

Condition	Items	Abst.	Cand.	Eval.
A (GitHub)	208	121	88	87
B (NHTSA)	142	13	131	128
Baseline (all SYRS)	350	73	277	274

Evaluation totals $87 + 128 + 274 = 489$ VCs spanning 307 unique SYRS; 43 items yielded no evaluated VC. Ch-A and Ch-B share no SYRS; the baseline spans all 350 items (overlapping both channels).

C. Evaluation Metrics

a) *Accept rate*: Fraction of VCs whose final label is Accept (three-criterion protocol: $C \geq 3$ AND $N = Y$ AND $U = Y$).

b) *Wilson 95% CI*: Confidence interval for the accept rate using the Wilson score method [31].

c) *Fisher’s exact test + Holm–Bonferroni correction*: Each channel is tested against the baseline with a one-sided Fisher’s exact test [32] for the directional RQ1 alternative that failure-grounded acceptance exceeds the no-failure-evidence baseline; p -values are corrected with Holm–Bonferroni [33] across three simultaneous comparisons (Ch-A, Ch-B, and A+B combined vs. baseline).

d) *Cohen’s h* : Effect size for proportion comparisons [34]: $h \geq 0.2$ small; $h \geq 0.5$ medium; $h \geq 0.8$ large.

e) *Generation-rate proxy*: The *generation rate* is the fraction of SYRS items for which the model generates a VC rather than abstaining (NO_NOVEL_VC_FOUND). Under active filtering a high rate indicates genuinely new failure angles, while near-universal generation ($>93\%$) signals the model is ignoring the abstention instruction. We use it as a lightweight quality proxy for the RQ4 sensitivity runs, where full human evaluation is unavailable.

f) *Abstention rates*: Ch-A abstains on 121 of 208 SYRS items (58.2%), Ch-B on 13 of 142 (9.2%). The gap reflects vocabulary alignment: once filtered, NHTSA fault modes map tightly to body-control and chassis items, whereas GitHub vocabulary is more varied relative to communication-stack and diagnostics items—the pipeline generates selectively, not indiscriminately.

g) *Retrieval precision*: To gauge upstream retrieval quality, a single author manually evaluated a stratified random sample of 200 requirement–evidence pairs (100 per channel, seed=42, similarity 0.45–0.78); the absence of a second annotator here is a construct-validity limitation. A pair is labeled *relevant* if the retrieved evidence describes a failure mode or boundary condition observable in the ECU under test (not merely a topically adjacent software defect or host-toolchain issue). The pipeline achieves an overall matching precision of **73.5 %** (Ch-A GitHub: **61.0 %**; Ch-B NHTSA: **86.0 %**), indicating that most retrieved evidence is failure-relevant, with Ch-A retrieval the weaker side. All prompts, model identifiers, run scripts, and raw LLM outputs for the reported runs are released with the replication package.

D. Inter-Annotator Agreement

a) *Rater profiles*: Four raters span the three ASPICE Verification & Validation (V&V) stakeholder roles (process ownership, academic methodology, hands-on practice), guarding against single-perspective bias: **E1** (Embedded Systems Project Manager, 10 yr) adjudicates split votes under defined criteria; **E2** (ASPICE assessor and functional-safety consultant) contributes SYS.2 process authority; **E3** (PhD candidate in automotive software testing) provides methodology-grounded judgment; **E4** (senior functional-test engineer, OEM/Tier-1) anchors physical V&V executability. This industry–academia panel follows CrUISE-AC’s use of complementary professional backgrounds [4].

b) *Main evaluation (RQ1/RQ2)*: E2, E3, and E4 independently scored the 489 VCs (each receiving two independent external ratings); E1 resolved split decisions under the three-criterion protocol, and the final label is the agreement of the two external raters or E1’s adjudication for splits. We use Gwet’s AC1 [35], [36] rather than Cohen’s κ to avoid the kappa paradox [37] under high Accept-class prevalence, reported on Completeness ($C \geq 3$).

c) *Ablation evaluation (RQ3)*: The 50-item subset comprises the first 50 Ch-A items in corpus order, all Communication-Stack requirements; this homogeneity strengthens internal comparability (category cannot confound the contrast) but restricts RQ3’s external scope to that category (§VI). All four raters (E1–E4) score the subset; the label is the majority vote of E2–E4 ($\geq 2/3$) with E1 as adjudicator (matching the main protocol), and all four scores feed the pairwise AC1 analysis.

E. Baselines

a) *No-failure-evidence baseline*: Qwen-Max prompted with the SYRS item and its existing golden VCs, but no retrieved failure context. This represents VC generation without external failure evidence and is the primary comparison point for RQ1.

b) *CrUISE-AC [4]*: The closest published system, targeting acceptance criteria for user stories from crowd-reported bugs. Accept rates and IAA values are compared in §V-A to contextualise our results.

V. RESULTS AND ANALYSIS

A. RQ1: Effectiveness of Failure-Grounded Prompting

Table III summarizes acceptance for all conditions; both channels significantly outperform the baseline after Holm–Bonferroni correction. Channel A (GitHub) reaches 65.5 % (57/87; OR=2.30, Cohen’s $h = 0.41$, Holm- $p = 0.0007$) and Channel B (NHTSA) 67.2 % (86/128; OR=2.48, $h = 0.45$, Holm- $p < 0.001$); combined, A+B reaches 66.5 % (143/215; $p < 0.001$, OR=2.40, $h = 0.43$), indicating that both channels can contribute under this channel–domain mapping.

TABLE III
RQ1: ACCEPT RATES VS. BASELINE (THREE-CRITERION PROTOCOL;
FINAL ADJUDICATED LABEL). 95 % CI IS THE WILSON SCORE INTERVAL
(PERCENTAGE POINTS).

Condition	VCs	Acc.	Rate	95 % CI	Holm- p
Baseline (no-failure)	274	124	45.3 %	[39.5, 51.2]	—
Channel A (GitHub)	87	57	65.5 %	[55.1, 74.7]	0.0007***
Channel B (NHTSA)	128	86	67.2 %	[58.7, 74.7]	<0.001***
A+B combined	215	143	66.5 %	[60.0, 72.5]	<0.001***

*** $p < 0.001$ (Holm–Bonferroni corrected).

a) *IAA*: Pairwise Gwet’s AC1 on Completeness ($C \geq 3$) is 0.84 (E2–E3), 0.65 (E2–E4), and 0.69 (E3–E4), averaging 0.73; the two lower pairs both involve E4’s physical-executability perspective vs. E2/E3’s requirements/process perspectives (a discipline-level finding, §VI). The mean is comparable to CrUISE-AC’s AC1 = 0.74 [4], under our more restrictive three-criterion acceptance predicate.

b) *Per-item accepted-VC yield*: Conditional accept rates compare *generated* VCs and are not commensurable across conditions with different abstention rates; the *per-item yield* (accepted VCs over all assigned items) gives a common denominator: baseline 79.1%/45.3%/0.35, Ch-A 41.8%/65.5%/0.27, and Ch-B 90.8%/67.2%/0.61 (generation/accept/yield). Thus Ch-A is a precision channel with lower coverage, whereas Ch-B covers most assigned items at a higher accept rate.

B. RQ2: Channel–Domain Fit

Because the two channels are mapped to disjoint requirement subdomains, RQ2 concerns where each channel fits rather than which channel performs better in a head-to-head comparison. Within its assigned domain, each channel yields a category-dependent uplift over the no-failure-evidence baseline averaging +20 percentage points (pp): +20.7 pp for Channel A and +21.3 pp for Channel B (category-matched,³ from Table IV), each backed by a small-to-medium effect size (Channel A $h = 0.41$; Channel B $h = 0.45$). The grounding mechanism differs by source: once Stage-2 filtering removes their colloquial noise, NHTSA complaints’ component-tagged fault modes align closely with SYRS vocabulary, whereas

³Pooled deltas (the +20/+22 pp in RQ1) compare each channel against the full baseline; *category-matched* deltas (RQ2) compare against the baseline restricted to that channel’s categories.

GitHub issues offer higher volume (696 vs. 364 records) at the cost of greater terminological variance.

a) *Scope-stratified acceptance*: Table IV breaks acceptance down by ECU category. Body-control and chassis items (Ch-B) show consistently high acceptance—IO/HMI 74.2 % (+28 pp), Functional Safety 71.9 %, Body Control 68.3 %—as filtered NHTSA narratives map tightly to sensor-level and state-machine specifications; communication-stack items (Ch-A) improve more modestly (61.3 % vs. 56.4 %, +4.9 pp), reflecting greater GitHub-issue vocabulary variation for protocol-specific SYRS items.

TABLE IV
SCOPE-STRATIFIED ACCEPTANCE BY ECU CATEGORY (THREE-CRITERION PROTOCOL, FINAL ADJUDICATED LABEL). DCM/DEM = DIAGNOSTIC COMMUNICATION/EVENT MANAGER. “—” = NO EVALUATED BASELINE VC FELL IN THIS CATEGORY.

Category	Ch.	Pipeline		Baseline	
		Acc/N	Rate	Acc/N	Rate
<i>Body-Control & Chassis (Ch-B)</i>					
IO / HMI	B	23/31	74.2 %	18/39	46.2 %
Functional Safety	B	23/32	71.9 %	—	—
Body Control	B	28/41	68.3 %	21/42	50.0 %
<i>Comm. Stack & State Mgmt (Ch-A/B)</i>					
Communication Stack	A	19/31	61.3 %	31/55	56.4 %
System State Mgmt	B	12/24	50.0 %	11/28	39.3 %
<i>Diagnostics (Ch-A)</i>					
Diagnostics (DCM)	A	29/44	65.9 %	25/64	39.1 %
Diagnostics (DEM)	A	4/6	66.7 %	14/33	42.4 %
Diagnostic Services	A	5/6	83.3 %	4/13	30.8 %

Table I traces an accepted Channel B VC end-to-end; with the Channel A running example of §III it illustrates the two channels’ complementary discovery roles.

C. RQ3: Ablation

Table V reports human acceptance rates on the 50-item Ch-A evaluation subset under the three ablation conditions.

TABLE V
RQ3: HUMAN ACCEPTANCE RATE ON 50-ITEM CH-A SUBSET (4-RATER MAJORITY VOTE, $C \geq 3$ AND $N=Y$ AND $U=Y$).

Condition	Accepted	Rate
Full Field2VC (FL, filter + embed)	29/50	58.0 %
Abl-LM (LLM-Match*)	29/50	58.0 %
Abl-NF (No Filter, embed)	41/50	82.0 %

*CrUISE-AC-style LLM binary relevance classification [4]; no triple filtering or multi-LLM ensemble.

a) *Embedding matching vs. LLM binary classification*: On the same 50-item subset, the two *unfiltered* ablations isolate the matching choice: Abl-NF (raw-text embedding) achieves 82.0 % versus 58.0 % for Abl-LM (single-LLM binary classifier, CrUISE-AC-style but without its ensemble [4]). The full pipeline (FL, filter + embed) also reaches 58.0 %; without a filter + LLM-match cell, this is an ablation pattern rather than a complete main-effect decomposition.

b) *Filtering as a corpus-scale selectivity/coverage control*: We do not dispute the local inversion: on these 50 items, removing the filter (Abl-NF, 82.0 %) scores higher than the full pipeline (58.0 %). That cost is invisible at this scale and surfaces only at corpus scale, where Abl-NF’s generation-rate proxy reaches 96–100 %: without the $\langle \text{Trigger} | \text{Protocol} | \text{Behavior} | \text{Fault Mode} \rangle$ triple, raw embedding matches thematically misaligned issues and the generator produces a VC for nearly every SYRS regardless of relevance. The full pipeline thus trades subset-level precision (58.0 % vs. 82.0 %) for suppression of this near-universal ungrounded generation—a precision/coverage choice we adopt to scale beyond the 50-item demonstration. We read this as a design trade-off, not a quality result: the volume the filter suppresses at corpus scale is not itself validated by human review—the proxy’s acknowledged limit, and a reason to use it only after checking filtering and sentinel adherence (§V-D).

c) *IAA (FL condition)*: On the 50-item FL evaluation, individual accept rates were E1 33/50 (66.0 %), E2 24/50 (48.0 %), E3 47/50 (94.0 %), E4 27/50 (54.0 %), for a majority vote ($E2+E3+E4 \geq 2/3$) of 29/50 (58.0 %). Six rows produced E1–MV disputes—five where E1 accepted but MV rejected (ExtAddr scope mismatch or tool-layer contamination), one (Row 013) where E1 rejected but MV accepted (BufferedReader ECU wake-up boundary)—with MV prevailing per protocol. Mean 4-rater Gwet’s AC1 is 0.63 for Abl-NF (substantial) and 0.37 for Abl-LM (moderate), corroborating the quality difference; pairwise agreement spans 66–92 % (Abl-NF) and 34–88 % (Abl-LM), the E3–E4 low driven by methodological differences on Extended-Addressing conditions (§VI).

D. RQ4: Sentinel-Adherence Sensitivity

TABLE VI
RQ4: SENTINEL-ADHERENCE SENSITIVITY. THE PRIMARY RUN IS HUMAN-EVALUATED (ACCEPT RATE); SENSITIVITY RUNS REPORT ONLY THE GENERATION-RATE PROXY. RUNS ARE INDEPENDENT RE-RUNS, SO QWEN-MAX’S CH-A RATE HERE (40.4 %) DIFFERS FROM THE MAIN RUN (41.8 %) BY NORMAL VARIANCE.

LLM	Ch-A	Ch-B
<i>Primary run — human accept rate</i>		
Qwen-Max	65.5 %	67.2 %
<i>Sensitivity runs — generation-rate proxy</i>		
<i>Sentinel-adherent</i>		
Qwen-Max (sensitivity)	40.4 %	93.7 %
Qwen-Plus	59.6 %	94.4 %
<i>Non-sentinel-adherent (>93 % on both channels)</i>		
DeepSeek-v4-Pro	99.0 %	100.0 %
DeepSeek-v4-Flash	100.0 %	99.3 %
Kimi-K2.6	94.2 %	97.2 %
GLM-5	98.6 %	98.6 %
Qwen3.6-Plus	100.0 %	99.3 %

a) *Sentinel adherence separates two models from the other five*: Five of seven models generate for >93 % of items regardless of knowledge-source relevance, effectively ignoring the NO_NOVEL_VC_FOUND sentinel. Only Qwen-Max and

Qwen-Plus emit it discriminatively, abstaining far more on Channel A’s varied GitHub vocabulary than on Channel B (Table VI). Sentinel adherence is thus itself a model-selection signal: the proxy holds only on instruction-adherent models, where controlled abstention matters as much as raw accuracy. These runs do not evaluate cross-model VC quality. On both sentinel-adherent models (the two Qwen variants) the *generation-rate* $B > A$ ordering holds; this is not the *accept-rate* ordering of RQ2, which, with only two adherent models, remains untested across families.

VI. DISCUSSION AND THREATS TO VALIDITY

A. Failure Mode Analysis

Analysis of the 72 rejected VCs across Channels A and B reveals three failure modes that account for the channel acceptance gap. All three are residual gaps the Stage-5 expert review exists to catch: Field2VC proposes candidate missing VCs for an engineer to ratify, not a finished SYS.2 baseline.

FM-1: Boundary Underspecification (the most common mode). The VC is well-formed but lacks quantitative parameters (timing constraints, byte constants, thresholds), because source issues describe qualitative fault behavior without numeric bounds and the LLM omits boundaries that appear only in the SYRS text. More common in Channel A; Channel B’s component metadata and filtered fault modes constrain the generated bounds.

FM-2: Scenario Transplant (dominant Channel-B mode). The VC identifies the right ECU component and fault type but embeds the test in an irrelevant consumer-driving scenario lifted from an NHTSA narrative (e.g., “drive... 30 minutes with slight steering inputs”), making it impossible on a standard Hardware-in-the-Loop (HIL) bench.

FM-3: Filtering Collapse (thematic misalignment at scale). Without Stage 2 filtering, embedding similarity alone retrieves topically adjacent but thematically misaligned issues, and the LLM generates coherent VCs testing conditions unrelated to the SYRS under test (e.g., a UDS service-0x29 test for a framing requirement). This mode surfaces *at corpus scale*, not on the 50-item subset: across the full corpus Abl-NF generates for 96–100 % of items (§V-C), whereas the filtering-enabled pipeline abstains on the majority of Ch-A items—the gap is the volume of misaligned matches that filtering suppresses.

B. Threats to Validity

a) Construct validity: The accept/reject label depends on raters applying the three-criterion protocol ($C \geq 3$ AND $N = Y$ AND $U = Y$). To mitigate subjectivity we recruited four raters—three independent scorers (E2–E4) and an internal adjudicator (E1)—with automotive RE backgrounds, and computed Gwet’s AC1 [35] to account for the kappa paradox [37] under high Accept-class prevalence. E1 was not blind to the generation channel, a confirmation-bias risk—particularly for Channel B (67.2 %), which relied heavily on E1’s adjudication. We anchored E1’s rulings to the pre-registered protocol; in

every one of the 66 adjudicated Channel-B acceptances, E2—an independently qualified ASPICE assessor—had also rated the item Accept, so E1 never overrode an E2 reject to manufacture an acceptance. Channel-A acceptances, however, do not depend on adjudication on these items: restricted to the *agreed subset* (items both external raters labeled identically), Channel A accepts 37/37 items versus 119/128 for the baseline, while Channel B accepts only 20/39 (51.3 %); the NHTSA gain therefore leans more heavily on adjudication. As a voluntary stress test we also report the headline rates under a conservative *no-E1 floor* (every split counted a reject; E1 resolved 50/87 Channel-A, 89/128 Channel-B, and 146/274 baseline splits): Channel A falls to 42.5 % against a 43.4 % baseline and Channel B to 15.6 %, so the per-channel uplift does not survive this worst case. Single-adjudicator, non-blind scoring thus remains a genuine limitation: unlike CrUISE-AC [4], which accepts on ≥ 3 -of-4 *independent* experts, our protocol pairs two external raters with an adjudicator, and a larger blinded, independent panel under majority vote is the natural strengthening. The two lowest Completeness pairs both involve E4, reflecting disciplinary divergence between E4’s executability perspective and E2/E3’s requirements/process view, not measurement failure. The generation-rate proxy (RQ4) approximates novelty without human evaluation but cannot rank novel VCs by quality; it should be interpreted only after verifying that the model honors the generation-refusal sentinel (§V-D). Final Accept agreement is lower than Completeness agreement because the three-criterion predicate compounds C/N/U disagreements; the adjudication footprint, agreed subset, and no-E1 floor quantify this sensitivity. We also do not verify that accepted VCs are conflict-free with VCs attached to other SYRS items; the protocol assesses local usefulness for the target requirement.

b) Evaluator methodology divergence (Abl-LM): The Abl-LM condition’s low E3–E4 agreement (34 %, AC1 = -0.24) is a systematic methodological difference, not noise: E3 judged Extended-Addressing preconditions on protocol-level merit (ISO 15765-2), whereas E2 and E4 treated them as deployment-blocking. Both are legitimate; majority vote lets the conservative position prevail, making the 58.0 % Abl-LM rate a lower bound.

c) Internal validity: Stage 1 (Issue Collection) uses repositories chosen for ECU-adjacent vocabulary; we mitigate by drawing from three distinct stacks (UDS diagnostics, CAN transport, ECU toolchain). A small prompt-budget asymmetry remains: the do-not-repeat golden-VC context is truncated at 800 characters in the pipeline conditions but 600 in the baseline; we disclose rather than re-run, as it affects only that context, not the failure evidence. The absence of a hard similarity gate (§III-C) is supported post-hoc: every VC surviving sentinel abstention had a top-1 similarity ≥ 0.50 , and above this floor similarity does not predict acceptance (quartile accept rates 57–76 %, non-monotonic); a hypothetical 0.60 gate would have discarded VCs with a *higher* accept rate (70.0 %) than those retained (63.5 %). The Channel-B trace (Table I) illustrates why: even when a retrieved complaint

diverges in mechanism from the requirement, requirement-grounded generation—conditioned on the SYRS and its golden VC, not the evidence alone—keeps the VC anchored to SYRS intent, provided the requirement states that intent precisely enough.

d) *Training-data exposure*: Public Channel-A issues are plausibly in the generator’s pretraining data, a memorisation concern. The no-failure-evidence baseline reduces, but does not eliminate, this concern: the same Qwen-Max without retrieved failure evidence accepts only 45.3 %, so any memorised knowledge is not applied as often without in-context evidence; the same pattern appears on Channel B.

e) *Evidence grounding vs. abstention*: One might ask whether the gains reflect the failure evidence or merely the Stage-2 abstention reporting only a favorable subset. The per-item yield (§V-A), which charges every abstention as a miss, separates the two: Channel B attains its +21.3pp uplift while abstaining on only 9.2 % of items, so it cannot be selective silence, whereas Channel A abstains on 58.2 % and its yield (0.27) falls below baseline (0.35), trading recall for precision. We therefore claim failure grounding improves VC quality where filtered evidence aligns with requirement vocabulary, not that it uniformly raises generation; an abstention-matched head-to-head is left to future work.

f) *DUT-boundary filtering failure (RQ3)*: In the 50-item FL evaluation, 14 of the 21 MV-Reject VCs (67 %) arose from Device-Under-Test (DUT) boundary leakage: host-layer properties of the test libraries (log levels, file-descriptor limits, periodic-sender timing) were extracted as ECU-observable because Stage 2 cannot infer which side of the DUT boundary a property belongs to. For SYRS.0693_A, the generator asserted correctness at file-descriptor counts beyond 1023—a `python-can` property, not ECU-observable—which all but one rater rejected, indicating that the evaluation is not rubber-stamping plausible output. A namespace guard for known test tools before triple extraction would suppress most such false positives.

g) *External validity*: The primary evaluation covers one proprietary project. To probe generalizability we applied Channel A (Qwen-Max) to a separate Domain-2 pilot of 85 Diagnostics SYRS items (DCM, DEM, Diagnostic Services), which produced 15 novel VCs (17.6 % generation-rate proxy). A three-rater evaluation (E2–E4) accepted **13 of 15** (86.7 %, Wilson 95 % CI 62–96 %), exceeding the primary Channel-A rate (65.5 %), with substantial Novelty agreement (AC1=0.69) and moderate Usefulness agreement (0.42, lowered by E4’s stricter executability criterion). The overlapping CI gives only preliminary evidence that cross-domain transfer succeeds without domain-specific prompting. Further limits: the B > A ordering holds only on two Qwen-family models, so cross-family stability is untested; powertrain, ADAS, and gateway domains remain unexamined; NHTSA effectiveness may not transfer where no regulator-operated complaint database exists; and the RQ3 ablation subset is homogeneous (all 50 items Communication-Stack).

h) *Conclusion validity*: We apply Holm–Bonferroni correction [33] across three Fisher’s exact tests (family-wise $\alpha = 0.05$), report Wilson 95 % CIs [31] for all rates, and report Cohen’s h [34] alongside p -values to curb over-interpretation in small-to-medium samples. Fisher’s exact test assumes independent observations; because VCs are nested within requirements and scored by a shared rater pool, that assumption holds only approximately, so we read the effect sizes and the sensitivity analyses as checks against the resulting clustering. Several scope-stratified cells (Table IV) are small (Diagnostics DEM and Diagnostic Services each $n=6$), with wide Wilson intervals; we report them as descriptive breakdowns, not powered comparisons.

VII. CONCLUSION

We presented Field2VC, a five-stage pipeline for completing automotive ECU verification criteria from empirical failure evidence. It grounds LLM generation in two complementary external channels operating at distinct abstraction levels—developer-reported GitHub defects (code level) and regulator-collected NHTSA complaints (vehicle level)—filtering each report into a structured fault representation and abstaining when no retrieved evidence aligns. Across a four-expert evaluation (three independent scorers and one internal adjudicator) of 489 generated VCs spanning 307 unique ASPICE SYS.2 requirements, failure grounding raised expert acceptance from a 45.3 % no-failure-evidence baseline to 65.5 %–67.2 % (Holm-corrected $p < 0.001$, $h = 0.41$ – 0.45). Corpus-scale experiments further showed that abstention materially affected precision: without it, ungrounded generation became nearly universal (96–100 %).

These findings are best interpreted as characterizing when and where failure grounding is effective rather than as a single headline acceptance rate. An adjudication-sensitivity analysis (§VI) showed the effect was uneven and partly adjudicator-dependent: within the adjudication-independent agreed subset ($n=37$), every GitHub-grounded VC was accepted, versus 93 % for the baseline, whereas the NHTSA channel leaned more heavily on adjudication. Beyond the headline rate, the results suggest two design heuristics. First, *a failure source appears useful where, after filtering, its vocabulary matches the quantitative, component-level language that safety-critical requirements demand*—regulator complaints satisfy this for body-control and chassis requirements, developer issues for the communication and diagnostics stack, and the two corpora should not be treated as interchangeable without checking this alignment. Second, *the generation-rate proxy for novelty should be used only after verifying that the model honors an explicit generation-refusal instruction*—a precondition five of our seven models violated—so sentinel adherence must be verified before the proxy is trusted.

A Domain 2 pilot on 85 Diagnostics SYRS (UDS/DTC), where 13 of 15 novel VCs (86.7 %) were accepted without domain-specific prompting, provides preliminary support for the first heuristic in a separate Diagnostics corpus. Collectively, these results answer the four research questions: fail-

ure grounding improves expert acceptance over a no-failure-evidence baseline (RQ1); the gain depends on alignment between a requirement's subdomain and its failure source's filtered vocabulary (RQ2); the two unfiltered matching variants show different acceptance rates, while filtering chiefly controls corpus-scale ungrounded generation through abstention rather than subset-level quality (RQ3); and the channel ordering holds across LLMs only for models that honor the generation-refusal sentinel, which most do not (RQ4). To support replication, we provide a reproducible benchmark (Data Availability): the 350-item SYRS corpus (307 evaluated) with expert-annotated golden VCs, all 489 generated VCs with per-rater labels and adjudication, the three-criterion protocol, the prompts, and the analysis scripts.

Future Work. The immediate strengthening is a blinded, independent majority-vote panel that turns the adjudication-dependent part of the NHTSA effect into a firm estimate. Longer term, accepted VCs should become machine-checkable predicates linked to source failures by a traceability matrix, and the vocabulary-alignment rule should become a predictive channel-selection model for other ECU subdomains and regulated systems.

REFERENCES

- [1] ISO 26262-4:2018: *Road Vehicles—Functional Safety—Part 4: Product Development at the System Level*, International Organization for Standardization Std., 2018.
- [2] B. Boehm and V. R. Basili, "Software defect reduction top 10 list," *Computer*, vol. 34, no. 1, pp. 135–137, 2001.
- [3] *Automotive SPICE Process Reference and Assessment Model, Version 4.0*, VDA Quality Management Center, 2023.
- [4] S. Schwedt and T. Ströder, "From Bugs to Benefits: Improving User Stories by Leveraging Crowd Knowledge with CrUISE-AC," in *47th IEEE/ACM International Conference on Software Engineering, ICSE 2025, Ottawa, ON, Canada, April 26 - May 6, 2025*. Ottawa, ON, Canada: IEEE, 2025, pp. 1385–1395.
- [5] T. Ward, "Areas of improvement for autonomous vehicles: A machine learning analysis of disengagement reports," in *2024 4th Interdisciplinary Conference on Electrics and Computer (INTCEC)*, 2024, pp. 1–6.
- [6] J. R. Palit, V. K. Kumarasamy, and O. A. Osman, "Data-Driven Analysis of Crash Patterns in SAE Level 2 and Level 4 Automated Vehicles Using K-means Clustering and Association Rule Mining," 2026.
- [7] A. Mavin, P. Wilkinson, A. Harwood, and M. Novak, "Easy approach to requirements syntax (EARS)," in *Proc. 17th IEEE Int. Requirements Engineering Conf. (RE)*, Atlanta, Georgia, USA, 2009, pp. 317–322.
- [8] D. Luitel, S. Hassani, and M. Sabetzadeh, "Improving requirements completeness: Automated assistance through large language models," *Requirements Engineering*, vol. 29, pp. 73–95, 2023.
- [9] X. Lian, S. Wang, H. Zou, F. Liu, J. Wu, and L. Zhang, "Incorporating verification standards for security requirements generation from functional specifications," *Proc. ACM Softw. Eng.*, vol. 2, no. FSE, pp. 1710–1732, 2025.
- [10] F. Niu, R. Pan, L. C. Briand, H. Hu, and K. Koravadi, "TVR: Automotive System Requirement Traceability Validation and Recovery Through Retrieval-Augmented Generation," *CoRR*, 2025.
- [11] A. Masoudifard, M. M. Sorond, M. Madadi, M. Sabokrou, and E. Habibi, "Leveraging Graph-RAG and Prompt Engineering to Enhance LLM-Based Automated Requirement Traceability and Compliance Checks," 2024.
- [12] N. Alturayef, I. Ahmad, and J. Hassine, "TraceLLM: Leveraging large language models with prompt engineering for enhanced requirements traceability," *Requirements Engineering*, vol. 31, no. 1, pp. 6–54, 2026.
- [13] G. Long, J. Gong, H. Fang, and T. Chen, "Learning Software Bug Reports: A Systematic Literature Review," *ACM Transactions on Software Engineering and Methodology*, vol. 35, no. 7, pp. 1–47, 2026.
- [14] J. P. Meher, S. Biswas, and R. Mall, "Deep learning-based software bug classification," *Information and Software Technology*, vol. 166, p. 107350, 2024.
- [15] Q. A. Do, S. R. Chekuri, and T. Bhowmik, "Automated support to capture creative requirements via requirements reuse," in *Reuse in the Big Data Era*, X. Peng, A. Ampatzoglou, and T. Bhowmik, Eds., Cham, 2019, vol. 11602, pp. 47–63.
- [16] J. T. Almonte, S. A. Boominathan, and N. Nascimento, "Automated non-functional requirements generation in software engineering with large language models: A comparative study," in *2025 IEEE International Conference on Collaborative Advances in Software and Computing (CASCON)*, 2025, pp. 474–483.
- [17] V. Kosciński, S. Hashemi, and M. Mirakhorli, "On-Demand Security Requirements Synthesis with Relational Generative Adversarial Networks," in *45th IEEE/ACM International Conference on Software Engineering, ICSE 2023, Melbourne, Australia, May 14–20, 2023*. Melbourne, Australia: IEEE, 2023, pp. 1609–1621.
- [18] F. Wang, C. Arora, Y. Liu, K. Huang, C. Tantithamthavorn, A. Aleti, D. Sambathkumar, and D. Lo, "Multi-Modal Requirements Data-based Acceptance Criteria Generation using LLMs," in *40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16–20, 2025*. Seoul, Korea, Republic of: IEEE, 2025, pp. 3334–3345.
- [19] M. Krishna, B. Gaur, A. Verma, and P. Jalote, "Using LLMs in Software Requirements Specifications: An Empirical Evaluation," in *2024 IEEE 32nd International Requirements Engineering Conference (RE)*, Reykjavik, Iceland, 2024, pp. 475–483.
- [20] Z. Nan, Z. Guo, K. Liu, and X. Xia, "Test Intention Guided LLM-Based Unit Test Generation," in *47th IEEE/ACM International Conference on Software Engineering, ICSE 2025, Ottawa, ON, Canada, April 26 - May 6, 2025*. Ottawa, ON, Canada: IEEE, 2025, pp. 1026–1038.
- [21] D. Jin, Z. Jin, X. Chen, and C. Wang, "MARE: Multi-Agents Collaboration Framework for Requirements Engineering," *CoRR*, 2024.
- [22] S. Kirchner and A. Knoll, "Generating automotive code: Large language models for software development and verification in safety-critical systems," in *2025 IEEE Intelligent Vehicles Symposium (IV)*, 2025, pp. 813–820.
- [23] D. Molinelli, L. Di Grazia, A. Martin-Lopez, M. D. Ernst, and M. Pezzè, "Do LLMs Generate Useful Test Oracles? An Empirical Study with an Unbiased Dataset," in *40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16–20, 2025*. Seoul, Korea, Republic of: IEEE, 2025, pp. 278–290.
- [24] N. Petrovic, V. Zolfaghari, F. Pan, and A. Knoll, "LLM-empowered functional safety and security by design in automotive systems," 2026.
- [25] B. V. Balu, F. Geissler, F. Carélla, J.-V. Zacchi, J. Jiru, N. Mata, and R. Stolle, "Towards automated safety requirements derivation using agent-based RAG," in *AAAI Spring Symposia*, 2025.
- [26] National Highway Traffic Safety Administration, "Office of defects investigation (ODI) complaints and investigations database," NHTSA, Tech. Rep., 2024, available at <https://api.nhtsa.gov/>.
- [27] C. D. Manning, P. Raghavan, and H. Schütze, *Introduction to Information Retrieval*, 1st ed., Cambridge University Press, 2008.
- [28] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proc. 2019 Conf. Empirical Methods in Natural Language Processing (EMNLP)*, Hong Kong, China, 2019, pp. 3980–3990.
- [29] S. Zhao, Y. Shao, Y. Huang, J. Song, Z. Wang, C. Wan, and L. Ma, "Understanding the Fundamental Design Decisions of Retrieval-Augmented Generation Systems," *ACM Transactions on Software Engineering and Methodology*, p. 3802824, 2026.
- [30] C. Gao, G. Fan, C. Y. Chong, S. Chen, C. Liu, D. Lo, Z. Zheng, and Q. Liao, "A Systematic Literature Review of Code Hallucinations in LLMs: Characterization, Mitigation Methods, Challenges, and Future Directions for Reliable AI," 2025.
- [31] E. B. Wilson, "Probable inference, the law of succession, and statistical inference," *Journal of the American Statistical Association*, vol. 22, no. 158, pp. 209–212, 1927.
- [32] R. A. Fisher, "On the interpretation of χ^2 from contingency tables, and the calculation of P," *Journal of the Royal Statistical Society*, vol. 85, no. 1, pp. 87–94, 1922.
- [33] S. Holm, "A simple sequentially rejective multiple test procedure," *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.
- [34] J. Cohen, *Statistical Power Analysis for the Behavioral Sciences*, 2nd ed. Hillsdale, NJ: Lawrence Erlbaum Associates, 1988.
- [35] K. L. Gwet, "Computing inter-rater reliability and its variance in the presence of high agreement," *British Journal of Mathematical and Statistical Psychology*, vol. 61, no. 1, pp. 29–48, 2008.
- [36] —, *Handbook of Inter-Rater Reliability: The Definitive Guide to Measuring the Extent of Agreement Among Raters*, 4th ed. Advanced Analytics, LLC, 2014.
- [37] A. R. Feinstein and D. V. Cicchetti, "High agreement but low Kappa: I. The problems of two paradoxes," *Journal of Clinical Epidemiology*, vol. 43, no. 6, pp. 543–549, 1990.